



DEVOXX
FRANCE

SparkSQL pour analyser vos données Cassandra

#TiaSparkSQL

@alexanderDeja @maxospiquante

Qui sommes-nous ?

DEVOXX
FRANCE

Alexander DEJANOVSKI

 @alexanderDeja

Développeur




chronopost

Maxence LECOINTE

 @maxospiquante

Développeur



chronopost

Cassandra

- Base NoSQL distribuée
- Langage de requête : CQL~=SQL
 - *SELECT * FROM ze_table WHERE ze_key=1*
- Pas de jointure, pas de group by, pas d'insert/select

Spark

- Map/Reduce en mémoire
- 10x-100x plus rapide que Hadoop
- **Scala, Java ou Python**
- Modules : Spark Streaming, MLlib, GraphX, **SparkSQL**

Objectif

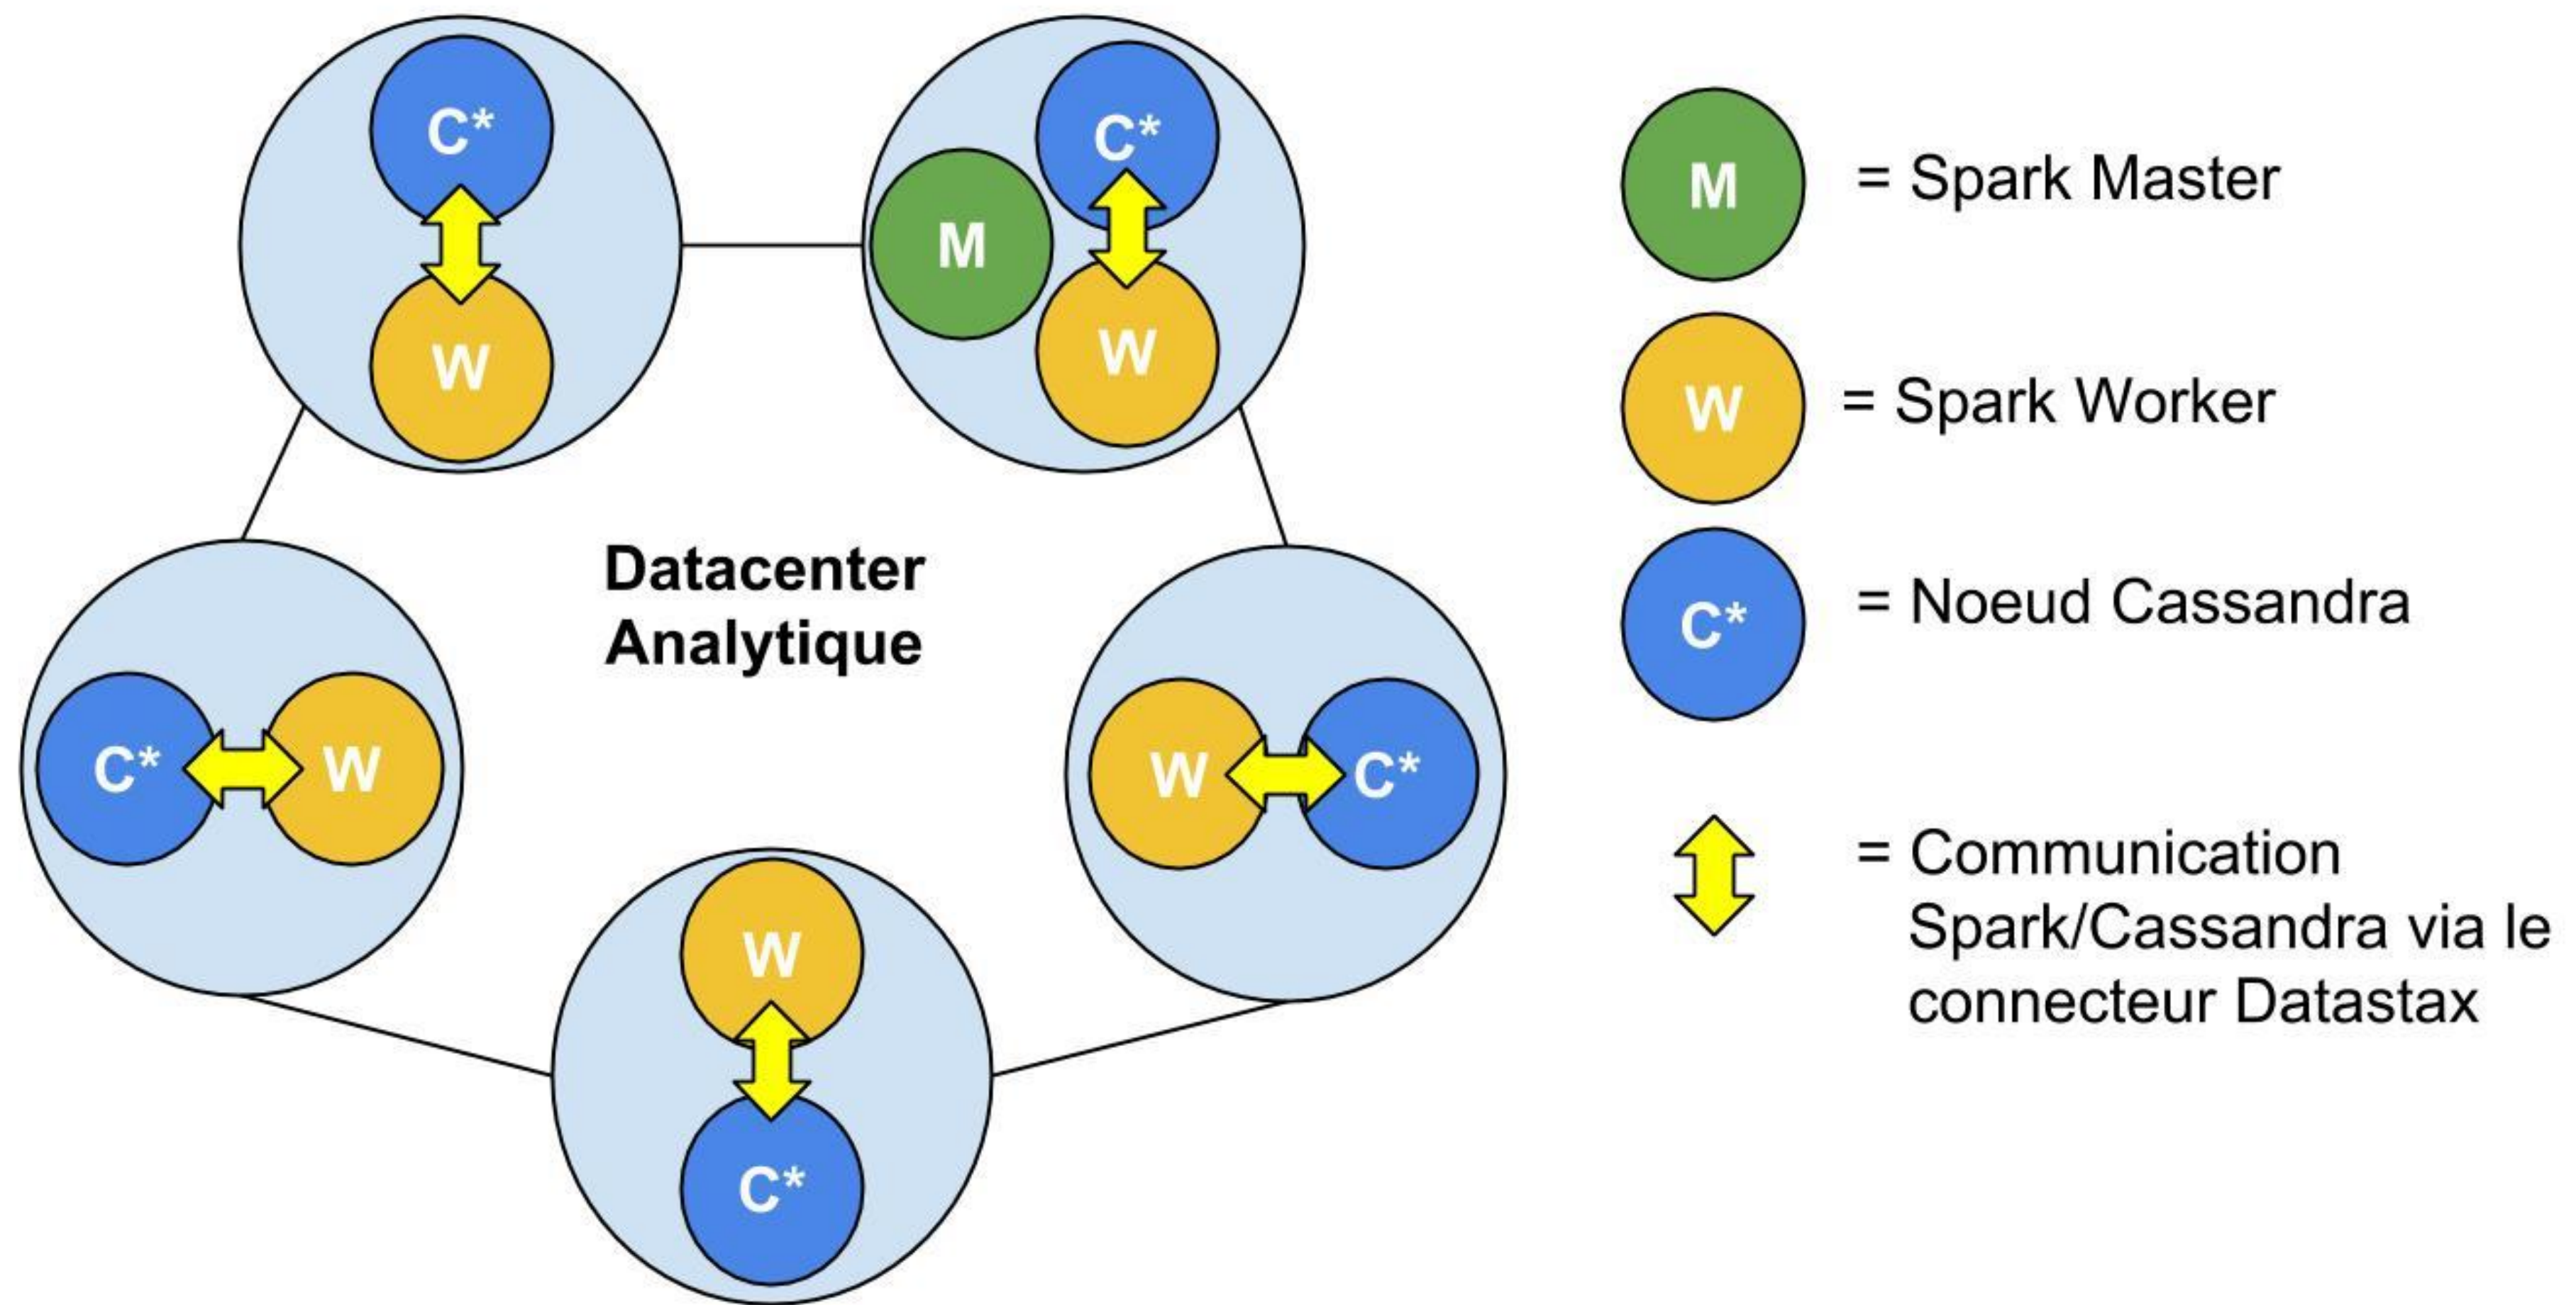
Cassandra << >> SparkSQL

Création de tables d'index

Calcul de statistiques (simples...)

sur les confs Devovx FR de 2012 à 2015

Datastax Spark Cassandra Connector



Setup

- Spark 1.1 ou 1.2 pour **Scala** et **Java**
- Connecteur Datastax :
<http://github.com/datastax/spark-cassandra-connector>
- Spark 1.1 pour **Python**
- Connecteur Calliope de TupleJump :
<http://tuplejump.github.io/calliope/start-with-sql.html>

Pour vous éviter (certaines) galères...

- Sources de ce TIA : <https://github.com/adejanovski/devoxx2015>
- Lisez le README



C'est quoi un RDD ?

- **Resilient Distributed Dataset**
- Collection d'objet distribuée et résiliente
- Permet le stockage de n'importe quel format de donnée

Schéma

Source

speaker		
id_speaker	text	K
societe	text	
nom_speaker	text	
twitter	text	

talk		
annee	int	K
titre	text	C
speakers	set<text>	
type_talk	text	

Schéma

Source

speaker		
id_speaker	text	K
societe	text	
nom_speaker	text	
twitter	text	

talk		
annee	int	K
titre	text	C
speakers	set<text>	
type_talk	text	

Statistiques

societe_par_annee		
societe	text	K
annee	int	C
nb	int	

speaker_par_annee		
nom_speaker	text	K
annee	int	C
nb	int	
id_speaker	text	

keyword_par_annee		
keyword	text	K
annee	int	C
nb	int	

Index

talk_par_speaker		
id_speaker	text	K
type_talk	text	C
titre	text	C
annee	text	

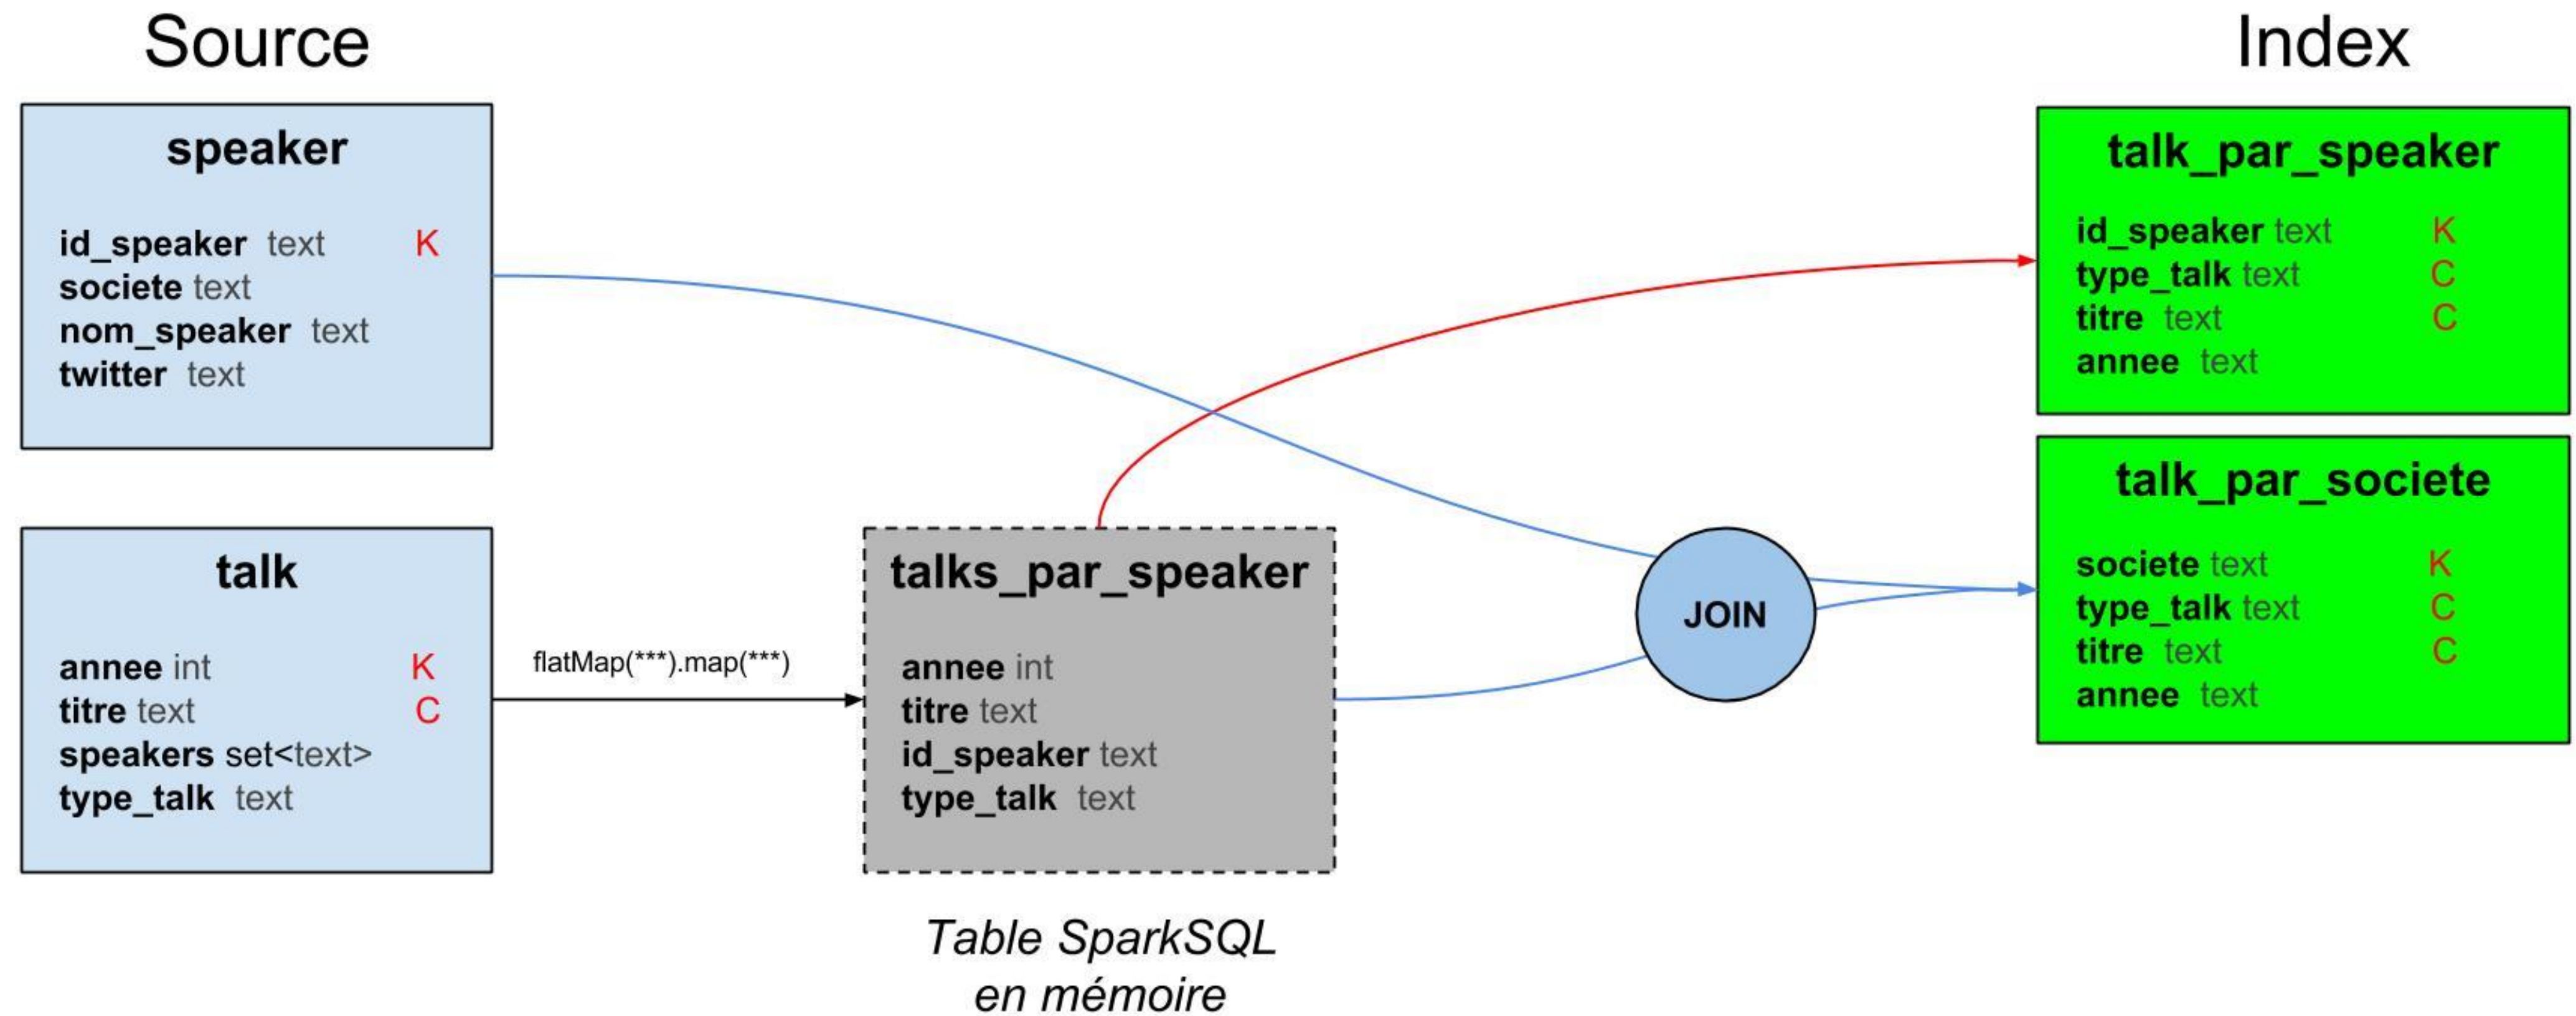
talk_par_societe		
societe	text	K
type_talk	text	C
titre	text	C
annee	text	

speaker_par_societe		
societe	text	K
id_speaker	text	C
nom_speaker	text	

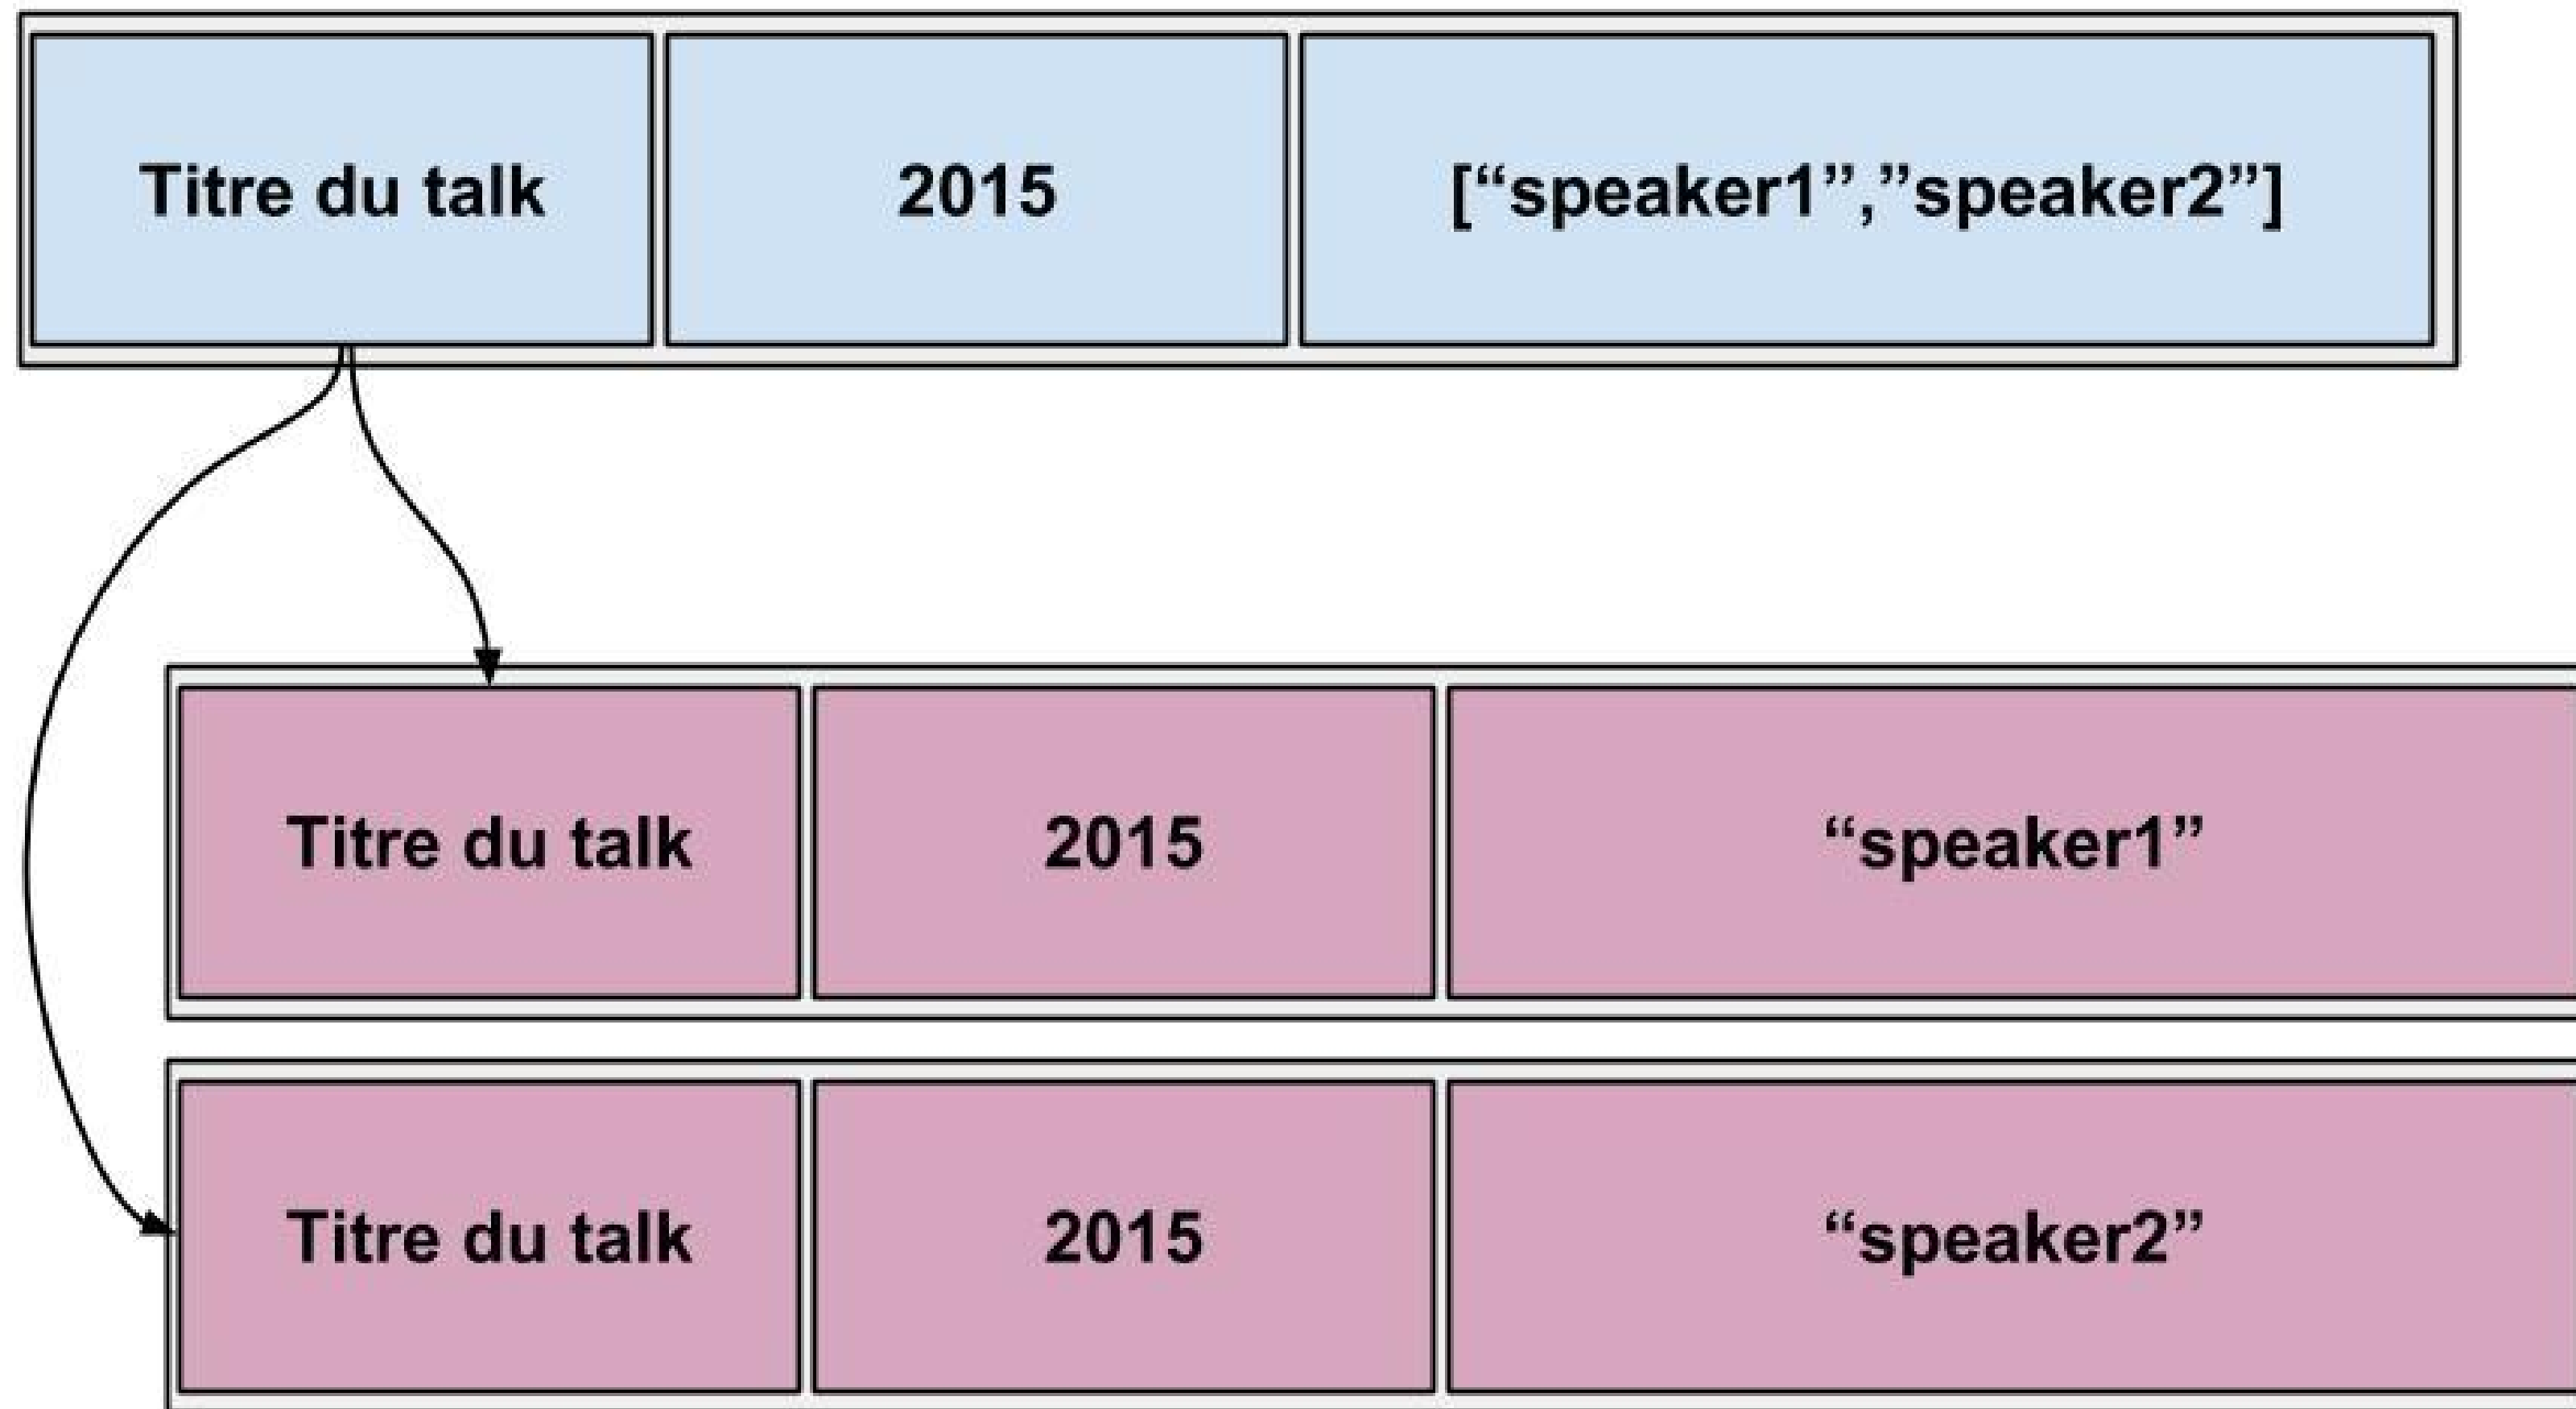
Etape 1



Scala-Fu



Scala-Fu : split par speaker



Code Scala

```
val rddTalk = cc.sql("select annee, titre, speakers, type_talk  
from devoxx.talk")  
  
// On sort de SparkSQL pour retravailler les données  
val splitBySpeakersRdd =  
rddTalk.flatMap(r => (r(2).asInstanceOf[scala.collection.immutable.Set[String]])  
  .map(m => (m,r) ))  
  
case class Talk(titre: String, speaker: String, annee: Int, type_talk: String)  
val talksSchemaRdd = splitBySpeakersRdd.map(  
  t => Talk(t._2.getString(1),t._1,t._2.getInt(0),t._2.getString(1),t._2.getString(3)))  
  
talksSchemaRdd.registerTempTable("talks_par_speaker")
```

Code Scala

```
val rddTalk = cc.sql("select annee, titre, speakers, type_talk  
from devoxx.talk")
```

// On sort de SparkSQL pour retravailler les données

```
val splitBySpeakersRdd =  
rddTalk.flatMap(r => (r(2).asInstanceOf[scala.collection.immutable.Set[String]])  
  .map(m => (m,r) ))
```

```
case class Talk(titre: String, speaker: String, annee: Int, type_talk: String)  
val talksSchemaRdd = splitBySpeakersRdd.map(  
  t => Talk(t._2.getString(1),t._1,t._2.getInt(0),t._2.getString(1),t._2.getString(3)))
```

```
talksSchemaRdd.registerTempTable("talks_par_speaker")
```

Code Scala

```
val rddTalk = cc.sql("select annee, titre, speakers, type_talk
                    from devovx.talk")

// On sort de SparkSQL pour retravailler les données
val splitBySpeakersRdd =
  rddTalk.flatMap(r => (r(2).asInstanceOf[scala.collection.immutable.Set[String]])
    .map(m => (m,r) ))

case class Talk(titre: String, speaker: String, annee: Int, type_talk: String)
val talksSchemaRdd = splitBySpeakersRdd.map(
  t => Talk(t._2.getString(1),t._1,t._2.getInt(0),t._2.getString(1),t._2.getString(3)))

talksSchemaRdd.registerTempTable("talks_par_speaker")
```


Code Scala

```
val rddTalk = cc.sql("select annee, titre, speakers, type_talk  
                    from devoxx.talk")  
  
// On sort de SparkSQL pour retravailler les données  
val splitBySpeakersRdd =  
  rddTalk.flatMap(r => (r(2).asInstanceOf[scala.collection.immutable.Set[String]])  
    .map(m => (m,r) ))  
  
case class Talk(titre: String, speaker: String, annee: Int, type_talk: String)  
val talksSchemaRdd = splitBySpeakersRdd.map(  
  t => Talk(t._2.getString(1),t._1,t._2.getInt(0),t._2.getString(1),t._2.getString(3)))  
  
talksSchemaRdd.registerTempTable("talks_par_speaker")
```

Code Scala : insertion Cassandra

```
cc.sql("insert into devoxx.talk_par_speaker  
      select speaker, type_talk, titre, annee  
      from talks_par_speaker").collect()
```

Code Scala : insertion Cassandra

```
val connector = CassandraConnector(sc.getConf)

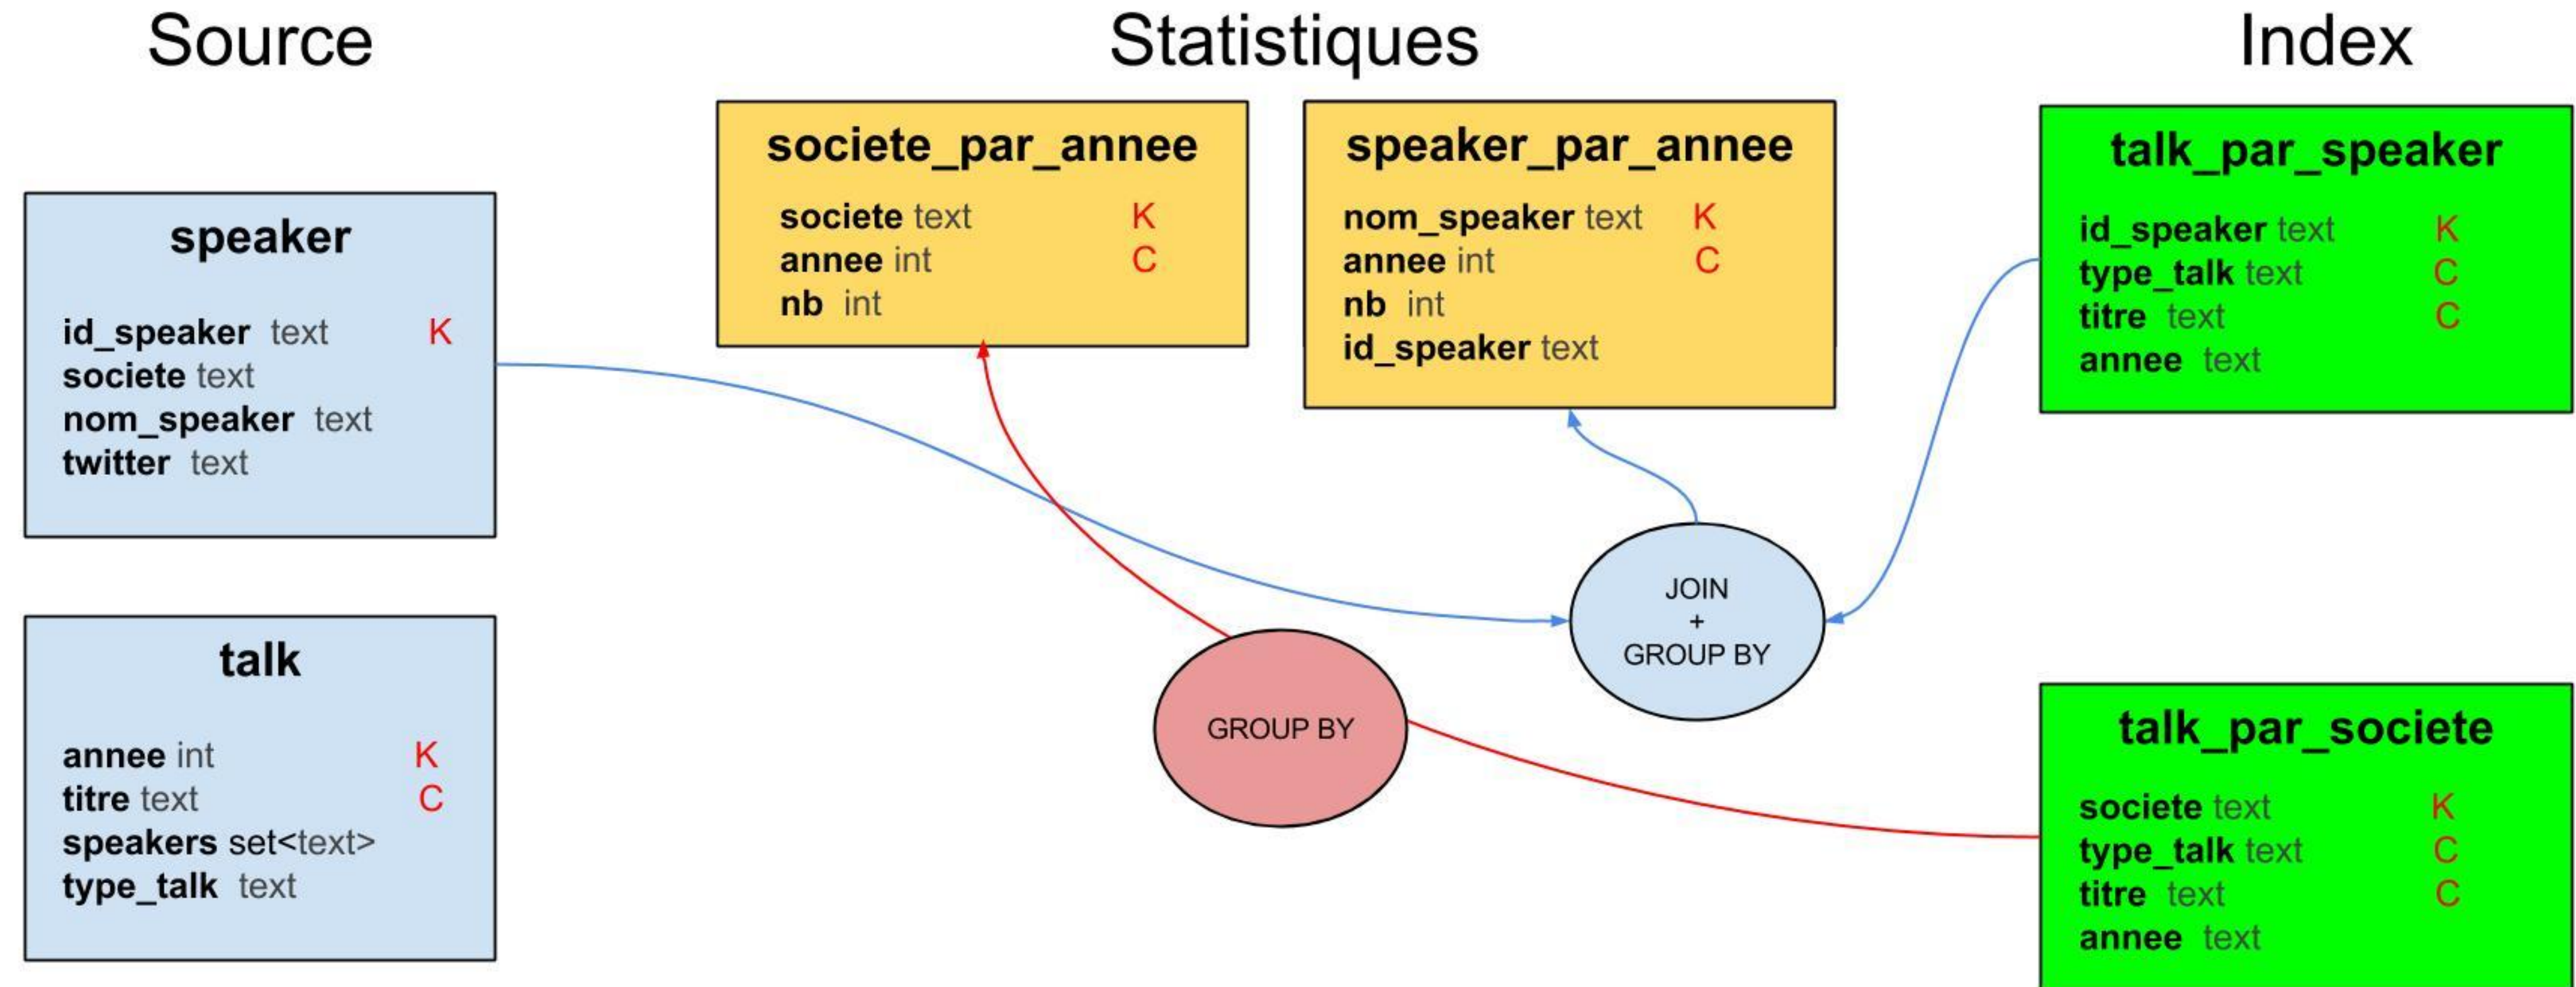
talksSchemaRdd.foreachPartition(partition => {
  connector.withSessionDo{ session =>
    partition.foreach(r => session.execute(
      "UDPATE devoxx.talk_par_speaker USING TTL ? " +
      set type_talk=?, titre=?, annee=? " +
      WHERE id_speaker = ?"),
      86400, r.type_talk, r.titre,
      r.annee.asInstanceOf[java.lang.Integer], r.speaker)
    })
})
```


“Demo time”

Etape 2



Java-Fu



Code Java

```
SchemaRDD nbTalkParSpeaker = cassandraSQLContext.sql(  
    "SELECT B.nom_speaker as nom_speaker, A.annee as annee,  
    "A.id_speaker as id_speaker " +  
    "FROM devoxx.talk_par_speaker A JOIN devoxx.speakers B "+  
    "ON A.id_speaker = B.id_speaker ");  
  
nbTalkParSpeaker.registerTempTable("tmp_talk_par_speaker");  
  
cassandraSQLContext.sql(  
    "INSERT INTO devoxx.speaker_par_annee " +  
    "SELECT nom_speaker, annee, count(*) as nb "+  
    "FROM tmp_talk_par_speaker group by nom_speaker, annee").collect();
```


Code Java

```
SchemaRDD nbTalkParSpeaker = cassandraSQLContext.sql(
    "SELECT B.nom_speaker as nom_speaker, A.annee as annee,
    "A.id_speaker as id_speaker " +
    "FROM devoxx.talk_par_speaker A JOIN  devoxx.speakers B "+
    "ON A.id_speaker = B.id_speaker ");

nbTalkParSpeaker.registerTempTable("tmp_talk_par_speaker");

cassandraSQLContext.sql(
    "INSERT INTO devoxx.speaker_par_annee " +
    "SELECT nom_speaker, annee, count(*) as nb "+
    "FROM tmp_talk_par_speaker group by nom_speaker, annee").collect();
```

Code Java

```
SchemaRDD nbTalkParSpeaker = cassandraSQLContext.sql(
    "SELECT B.nom_speaker as nom_speaker, A.annee as annee,
    "A.id_speaker as id_speaker " +
    "FROM devoxx.talk_par_speaker A JOIN devoxx.speakers B "+
    "ON A.id_speaker = B.id_speaker ");

nbTalkParSpeaker.registerTempTable("tmp_talk_par_speaker");

cassandraSQLContext.sql(
    "INSERT INTO devoxx.speaker_par_annee " +
    "SELECT nom_speaker, annee, count(*) as nb, id_speaker "+
    "FROM tmp_talk_par_speaker "+
    "GROUP BY nom_speaker, annee, id_speaker").collect();
```

Submit Java

```
./spark-submit  
  --class devoxx.Devoxx.....  
  --master spark://127.0.0.1:7077  
  devoxxSparkSql.jar
```

\\

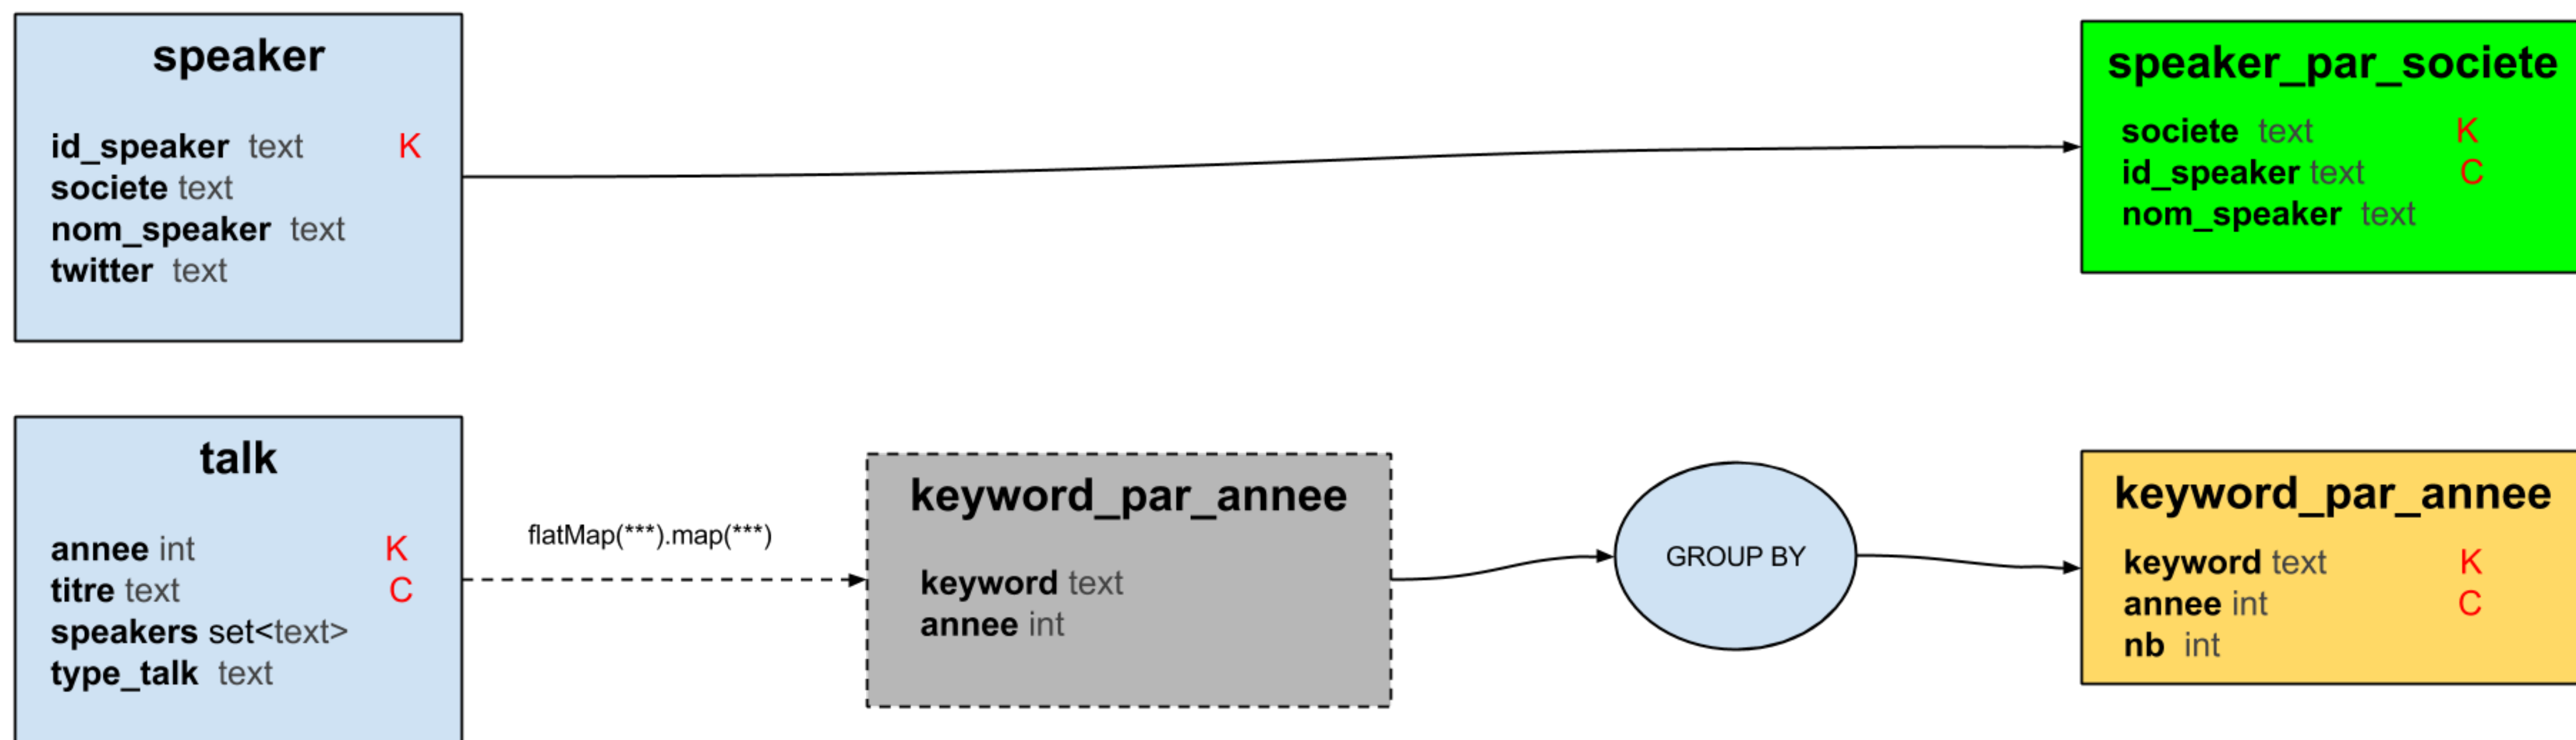
“Demo time”

Etape 3



Python-Fu

Source



Code Python

```
def split_keywords(row):  
    ## fonction splittant les titres par mot  
  
rddTalk = sqlContext.sql("SELECT titre, speakers, annee, categorie, type_talk FROM devoxx.talk")  
splitByKeywordRdd = rddTalk.flatMap(lambda r:split_keywords(r))  
splitByKeywordRdd_schema = sqlContext.inferSchema(  
    splitByKeywordRdd.filter(lambda word:len(word[0])>1)  
                        .map(lambda x:Row(keyword=x[0],annee=x[1])))  
splitByKeywordRdd_schema.registerTempTable("tmp_keywords")  
keyword_count = sqlContext.sql("""SELECT keyword, annee, count(*) as nb  
    FROM tmp_keywords  
    GROUP BY keyword, annee""")  
keyword_count_schema = sqlContext.inferSchema(keyword_count.map(lambda x:Row(...)))  
keyword_count_schema.registerTempTable("tmp_keywords_count")  
sqlContext.sql("""INSERT INTO devoxx.keyword_par_annee SELECT keyword, annee, nb  
    FROM tmp_keywords_count""")
```

Code Python

```
def split_keywords(row):
    ## fonction splittant les titres par mot

rddTalk = sqlContext.sql("select titre, speakers, annee, categorie, type_talk from devoxx.talk")
splitByKeywordRdd = rddTalk.flatMap(lambda r:split_keywords(r))
splitByKeywordRdd_schema = sqlContext.inferSchema(
    splitByKeywordRdd.filter(lambda word:len(word[0])>1)
                        .map(lambda x:Row(keyword=x[0],annee=x[1])))
splitByKeywordRdd_schema.registerTempTable("tmp_keywords")
keyword_count = sqlContext.sql("""SELECT keyword, annee, count(*) as nb
                                FROM tmp_keywords
                                GROUP BY keyword, annee""")
keyword_count_schema = sqlContext.inferSchema(keyword_count.map(lambda x:Row(...)))
keyword_count_schema.registerTempTable("tmp_keywords_count")
sqlContext.sql("""INSERT INTO devoxx.keyword_par_annee SELECT keyword, annee, nb
                FROM tmp_keywords_count""")
```


Code Python

```
def split_keywords(row):
    ## fonction splittant les titres par mot

rddTalk = sqlContext.sql("select titre, speakers, annee, categorie, type_talk from devoxx.talk")
splitByKeywordRdd = rddTalk.flatMap(lambda r:split_keywords(r))
splitByKeywordRdd_schema = sqlContext.inferSchema(
    splitByKeywordRdd.filter(lambda word:len(word[0])>1)
                        .map(lambda x:Row(keyword=x[0],annee=x[1])))
splitByKeywordRdd_schema.registerTempTable("tmp_keywords")
keyword_count = sqlContext.sql("""SELECT keyword, annee, count(*) as nb
                                FROM tmp_keywords
                                GROUP BY keyword, annee""")
keyword_count_schema = sqlContext.inferSchema(keyword_count.map(lambda x:Row(...)))
keyword_count_schema.registerTempTable("tmp_keywords_count")
sqlContext.sql("""INSERT INTO devoxx.keyword_par_annee SELECT keyword, annee, nb
                FROM tmp_keywords_count""")
```

Code Python

```
def split_keywords(row):
    ## fonction splittant les titres par mot

rddTalk = sqlContext.sql("select titre, speakers, annee, categorie, type_talk from devoxx.talk")
splitByKeywordRdd = rddTalk.flatMap(lambda r:split_keywords(r))
splitByKeywordRdd_schema = sqlContext.inferSchema(
    splitByKeywordRdd.filter(lambda word:len(word[0])>1)
                        .map(lambda x:Row(keyword=x[0],annee=x[1])))
splitByKeywordRdd_schema.registerTempTable("tmp_keywords")
keyword_count = sqlContext.sql("""SELECT keyword, annee, count(*) as nb
                                FROM tmp_keywords
                                GROUP BY keyword, annee""")
keyword_count_schema = sqlContext.inferSchema(keyword_count.map(lambda x:Row(...)))
keyword_count_schema.registerTempTable("tmp_keywords_count")
sqlContext.sql("""INSERT INTO devoxx.keyword_par_annee SELECT keyword, annee, nb
                FROM tmp_keywords_count""")
```

“Demo time”

Et voilà ! Des questions ?
